

Kcov - a single-step code coverage tool

Simon Kågström

Net Insight

<https://github.com/SimonKagstrom/kcov>

September 20, 2018

Motivation

```

    if ((err = ReadyHash(&SSLHashMD5, &hashCtx)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &clientRandom)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashMD5.final(&hashCtx, &hashOut)) != 0)
        goto fail;
}
else {
    /* DSA, ECDSA - just use the SHA1 hash */
    dataToSign = &hashes[SSL_MD5_DIGEST_LEN];
    dataToSignLen = SSL_SHA1_DIGEST_LEN;
}

hashOut.data = hashes + SSL_MD5_DIGEST_LEN;
hashOut.length = SSL_SHA1_DIGEST_LEN;
if ((err = SSLFreeBuffer(&hashCtx)) != 0)
    goto fail;

if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
    goto fail;
if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
    goto fail;

err = sslRawVerify(ctx,
                  ctx->peerPubKey,
                  dataToSign,          /* plaintext */
                  dataToSignLen,       /* plaintext length */
                  signature,
                  signatureLen);

if(err) {
    sslErrorLog("SSLDecodeSignedServerKeyExchange: sslRawVerify "
               "returned %d\n", (int)err);
    goto fail;
}
}

```

fail:

Motivation

```

90 0 / 1 if ((err = ReadyHash(&SSLHashMD5, &hashCtx)) != 0)
98 0 / 1 goto fail;
99 0 / 1 if ((err = SSLHashMD5.update(&hashCtx, &clientRandom)) != 0)
100 0 / 1 goto fail;
101 0 / 1 if ((err = SSLHashMD5.update(&hashCtx, &serverRandom)) != 0)
102 0 / 1 goto fail;
103 0 / 1 if ((err = SSLHashMD5.update(&hashCtx, &signedParams)) != 0)
104 0 / 1 goto fail;
105 0 / 1 if ((err = SSLHashMD5.final(&hashCtx, &hashOut)) != 0)
106 0 / 1 goto fail;
107 }
108 else {
109 /* DSA, ECDSA - just use the SHA1 hash */
110 1 / 1 dataToSign = &hashes[SSL_MD5_DIGEST_LEN];
111 1 / 1 dataToSignLen = SSL_SHA1_DIGEST_LEN;
112 }
113
114 1 / 1 hashOut.data = hashes + SSL_MD5_DIGEST_LEN;
115 1 / 1 hashOut.length = SSL_SHA1_DIGEST_LEN;
116 1 / 1 if ((err = SSLFreeBuffer(&hashCtx)) != 0)
117 0 / 1 goto fail;
118
119 1 / 1 if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
120 0 / 1 goto fail;
121 1 / 1 if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
122 0 / 1 goto fail;
123 1 / 1 if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
124 0 / 1 goto fail;
125 1 / 1 if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
126 goto fail;
127 1 / 1 goto fail;
128 if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
129 goto fail;
130
131 err = sslRawVerify(ctx,
132 ctx->peerPubKey,
133 dataToSign, /* plaintext */
134 dataToSignLen, /* plaintext length */
135 signature,
136 signatureLen);
137 if(err) {
138 sslErrorLog("SSLDecodeSignedServerKeyExchange: sslRawVerify "
139 "returned %d\n", (int)err);
140 goto fail;
141 }
142
143 fail:

```

More motivation

- Programming is hard!
- To improve software quality, good tooling and methodology is important



Outline

1 Overview

- Some words about code coverage
- Problems with traditional tools

2 Kcov usage

- Kcov overview
- Main features of kcov
- Integration with CI systems
- Python/Bash

3 How kcov works

- Elves, dwarves and breakpoints
- Implementation quirks on different architectures
- Implementation quirks and bugs
- Python and Bash coverage collection
- The design of kcov

Table of Contents

- 1 Overview
 - Some words about code coverage
 - Problems with traditional tools
- 2 Kcov usage
 - Kcov overview
 - Main features of kcov
 - Integration with CI systems
 - Python/Bash
- 3 How kcov works
 - Elves, dwarves and breakpoints
 - Implementation quirks on different architectures
 - Implementation quirks and bugs
 - Python and Bash coverage collection
 - The design of kcov

Code coverage terminology

- **Function coverage:** Functions in the program
- **Statement coverage:** Statements in the program
- **Branch coverage:** All ways through branches (if/case)
- **Condition coverage:** All boolean expressions evaluated to both true/false
- **Modified condition/decision coverage (MC/DC):** Each decision takes every possible outcome

(From Wikipedia)

Problems with traditional tools

- gcov + lcov is a multi-step process
- gcov leaves droppings after compilation/running
- A program which crashes will not generate coverage data

Example

```
$ gcc -g -Wall --coverage goto-fail.c
$ ./a.out
$ ls
a.out goto-fail.c goto-fail.gcda goto-fail.gcno
$ lcov --capture --directory project-dir --output-file coverage.info
$ genhtml coverage.info --output-directory out
```


Instead...



Table of Contents

1 Overview

- Some words about code coverage
- Problems with traditional tools

2 Kcov usage

- Kcov overview
- Main features of kcov
- Integration with CI systems
- Python/Bash

3 How kcov works

- Elves, dwarves and breakpoints
- Implementation quirks on different architectures
- Implementation quirks and bugs
- Python and Bash coverage collection
- The design of kcov

Kcov overview

- Kcov started as a fork of Bcov by Thomas Neumann in 2010
- Bcov doesn't rely on compile-time instrumentation, but instead uses debug information
- Bcov was a great idea, but I thought it could be improved upon

Kcov overview

- Kcov started as a fork of Bcov by Thomas Neumann in 2010
- Bcov doesn't rely on compile-time instrumentation, but instead uses debug information
- Bcov was a great idea, but I thought it could be improved upon
- Can you guess why it's called kcov?

Kcov overview

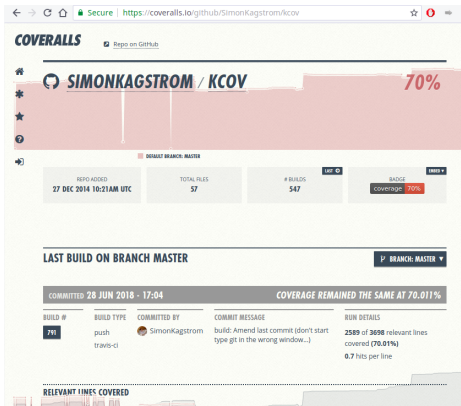
- Kcov started as a fork of Bcov by Thomas Neumann in 2010
- Bcov doesn't rely on compile-time instrumentation, but instead uses debug information
- Bcov was a great idea, but I thought it could be improved upon
- Can you guess why it's called kcov?
- It is not related to the kernel Kcov (and predates it by many years)

Main features of kcov

- Supports collection and reporting of binaries as long as debug information is present
- Automatically collects coverage from shared libraries
 - Both linked into the binary, and opened via **dlopen**
- Collection and reporting is done in a single step
- Generates HTML output as well as several XML formats for integration in other environments
- Accumulates coverage information between runs
- Automatically merges multiple binaries into a combined report
 - Merging of multiple reports can also be done by the tool
- Works on Linux, FreeBSD and Mac OSX on x86, ARM and PowerPC

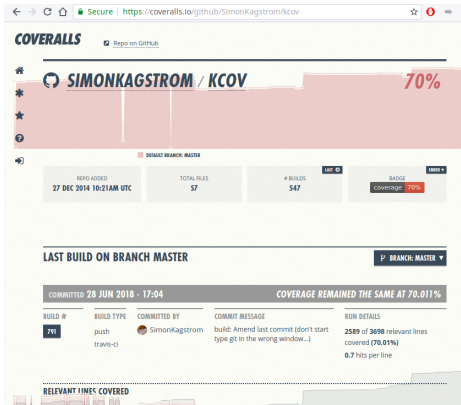
Integration with CI systems

- Jenkins and SonarQube output is generated by kcov
- Uploading to Coveralls.io is built-in
- Uploading to Codecov.io is supported by the upstream project



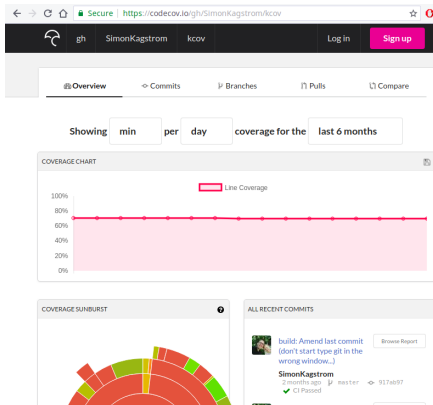
Integration with CI systems

- Jenkins and SonarQube output is generated by kcov
- **Uploading to Coveralls.io is built-in**
- Uploading to Codecov.io is supported by the upstream project



Integration with CI systems

- Jenkins and SonarQube output is generated by kcov
- Uploading to Coveralls.io is built-in
- Uploading to Codecov.io is supported by the upstream project



Python and Bash code coverage

- Kcov can also collect coverage for Python and Bash
- For Bash and Python, line counts are reported

```
Command: other.sh
Date: 2018-09-17 20:59:01
Code covered: 85.0%
Instrumented lines: 20
Executed lines: 17
```

```
1 #!/bin/sh
2
3 echo "This is another shell script"
4
5 if [ $1 -eq 5 ]; then
6     echo "Kalle anka"
7     fi
8
9 for var in "$@"; do
10     echo $var
11 done
12
13 round=0
14 22 for ( ; ; ) do
15 11 round=$((round + 1))
16 11 if [ $round -gt 10 ]; then
17 1 break
18 fi
19 10 echo "hej $round"
20 done
21
22 echo "Not covered" # LCOV_EXCL_LINE
23 1 echo "Covered"
24
25 # Ranges
26 echo "Not covered" # LCOV_EXCL_START
27 echo "Also not covered"
28 echo "Also not covered?" # LCOV_EXCL_STOP
29 1 echo "Covered"
30
```

```
Command: testdriver
Date: 2018-09-17 21:00:06
Code covered: 88.9%
Cov
```

```
1 #!/usr/bin/python
2
3 import sys
4 import unittest
5
6 class testcase(unittest.TestCase):
7
8     def setUp(self):
9         print('in setUp()')
10         self unittest = unittest.TestCase()
11
12     def testMethod1(self):
13         print('in testcase.testMethod1()')
14         expected=4
15         result= self unittest.testMethod1(2,2)
16
17         self.assertEqual(expected,result)
18
19
20 if __name__ == '__main__':
21     sys.path.append('.')
22
23     try:
24         import unittest
25     except:
26         print('exception')
27
28     unittest.main()
```

Interactive Demo!




Table of Contents

1 Overview

- Some words about code coverage
- Problems with traditional tools

2 Kcov usage

- Kcov overview
- Main features of kcov
- Integration with CI systems
- Python/Bash

3 How kcov works

- Elves, dwarves and breakpoints
- Implementation quirks on different architectures
- Implementation quirks and bugs
- Python and Bash coverage collection
- The design of kcov

Dwarf stabs Mach-O Elf

- **ELF** is a binary format used on many Unices
 - Both for object files, shared libraries and executable files
 - Contains the code, data, constants and relocation information
- **Mach-O** is the binary format used on Mac OS X
- **Dwarf** is a format for debug information
 - Contains the mapping between source lines and addresses
 - Type information for variables etc

Instrumenting binaries with kcov

- DWARF contains file:line to address records, which kcov uses to set breakpoints

Example

```
124:      return (void*)op_create(RPAR);
400ff7:      bf 30 00 00 00      mov     $0x30,%edi
400ffc:      e8 1f 05 00 00      callq  401520 <op_create>
401001:      e9 20 ff ff ff      jmpq   400f26 <ts_next_token+0x46>
401006:      66 2e 0f 1f 84 00 00  nopw   %cs:0x0(%rax,%rax,1)
40100d:      00 00 00
125: p_state->last_token_is_od = 0;
401010:      c7 43 08 00 00 00 00  movl   $0x0,0x8(%rbx)
126: return (void*)op_create(op);
401017:      bf 04 00 00 00      mov     $0x4,%edi
40101c:      e8 ff 04 00 00      callq  401520 <op_create>
```

Instrumenting binaries with kcov

- DWARF contains file:line to address records, which kcov uses to set breakpoints

Example

```
124:      return (void*)op_create(RPAR);
400ff7:      bf 30 00 00 00      mov     $0x30,%edi
400ffc:      e8 1f 05 00 00      callq  401520 <op_create>
401001:      e9 20 ff ff ff      jmpq   400f26 <ts_next_token+0x46>
401006:      66 2e 0f 1f 84 00 00  nopw   %cs:0x0(%rax,%rax,1)
40100d:      00 00 00
125: p_state->last_token_is_od = 0;
401010:      c7 43 08 00 00 00 00  movl    $0x0,0x8(%rbx)
126: return (void*)op_create(op);
401017:      bf 04 00 00 00      mov     $0x4,%edi
40101c:      e8 ff 04 00 00      callq  401520 <op_create>
```

Instrumenting binaries with kcov

- DWARF contains file:line to address records, which kcov uses to set breakpoints
- So how are breakpoints set on Linux?

Example

```
124:      return (void*)op_create(RPAR);
400ff7:      bf 30 00 00 00      mov     $0x30,%edi
400ffc:      e8 1f 05 00 00      callq  401520 <op_create>
401001:      e9 20 ff ff ff      jmpq   400f26 <ts_next_token+0x46>
401006:      66 2e 0f 1f 84 00 00  nopw   %cs:0x0(%rax,%rax,1)
40100d:      00 00 00
125: p_state->last_token_is_od = 0;
401010:      c7 43 08 00 00 00 00  movl   $0x0,0x8(%rbx)
126: return (void*)op_create(op);
401017:      bf 04 00 00 00      mov     $0x4,%edi
40101c:      e8 ff 04 00 00      callq  401520 <op_create>
```


How does kcov know where to set the breakpoint, really?

- The ELF binary contains program code in sections, located at address found in the binary
- But shared libraries and PIEs complicate this

Example

```
$ readelf --wide -S ~/local/bin/calc
```

There are 36 section headers, starting at offset 0x7ec8:

Section Headers:

[Nr]	Name	Type	Address	Off	Size
[...]					
[10]	.rela.plt	RELA	0000000000400a88	000a88	000150
[11]	.init	PROGBITS	0000000000400bd8	000bd8	000017
[12]	.plt	PROGBITS	0000000000400bf0	000bf0	0000f0
[13]	.text	PROGBITS	0000000000400ce0	000ce0	001b42
[14]	.fini	PROGBITS	0000000000402824	002824	000009
[15]	.rodata	PROGBITS	0000000000402830	002830	0006f5
[...]					

kcov on Linux, FreeBSD and Mac OSX

- ptrace is a truly archaic interface, and pretty non-portable
- libelf is also interesting: **elf_version** must be called at program start!
- Mac OSX has been implemented using the LLDB debugger as a library

Example

```
static long getRegs(pid_t pid, void *addr, void *regs, size_t len)
{
    #if defined(__aarch64__)
        struct iovec iov = {regs, len};
        return ptrace(PTRACE_GETREGSET, pid, (void *)NT_PRSTATUS, &iov);
    #else
        return ptrace((__ptrace_request) PTRACE_GETREGS, pid, NULL, regs);
    #endif
}
```

So are there any interesting bugs?



Dwarf quirks

- Dwarf generation on Linux is sometimes buggy, containing invalid entries

Example

```
124:      return (void*)op_create(RPAR);
400ff7:      bf 30 00 00 00      mov     $0x30,%edi
400ffc:      e8 1f 05 00 00      callq   401520 <op_create>
401001:      e9 20 ff ff ff      jmpq    400f26 <ts_next_token+0x46>
401006:      66 2e 0f 1f 84 00 00  nopw    %cs:0x0(%rax,%rax,1)
40100d:      00 00 00
125: p_state->last_token_is_od = 0;
401010:      c7 43 08 00 00 00 00  movl    $0x0,0x8(%rbx)
126: return (void*)op_create(op);
401017:      bf 04 00 00 00      mov     $0x4,%edi
40101c:      e8 ff 04 00 00      callq   401520 <op_create>
```

Dwarf quirks

- Dwarf generation on Linux is sometimes buggy, containing invalid entries

Example

```
124:      return (void*)op_create(RPAR);
400ff7:      bf 30 00 00 00      mov     $0x30,%edi
400ffc:      e8 1f 05 00 00      callq  401520 <op_create>
401001:      e9 20 ff ff ff      jmpq   400f26 <ts_next_token+0x46>
401006:      66 2e 0f 1f 84 00 00 nopl    %cs:0x0(%rax,%rax,1)
40100d:      00 00 00
125: p_state->last_token_is_od = 0;
401010:      c7 43 08 00 00 00 00 movl    $0x0,0x8(%rbx)
126: return (void*)op_create(op);
401017:      bf 04 00 00 00      mov     $0x4,%edi
40101c:      e8 ff 04 00 00      callq  401520 <op_create>
```

LLDB and fork

- The OSX port has another interesting quirk: forks are unsafe
- LLDB doesn't catch breakpoints from children, and will therefore crash in the child after a fork

Example

```
auto child = fork();
if (child < 0)
    printf("Error\n");
else if (child == 0)
    printf("In child\n"); // Breakpoint here is bad news!
else
{
    printf("In parent, child %d\n", child);
    wait(&status);
}
```

LLDB and fork

- The OSX port has another interesting quirk: forks are unsafe
- LLDB doesn't catch breakpoints from children, and will therefore crash in the child after a fork
- How can we workaround that?

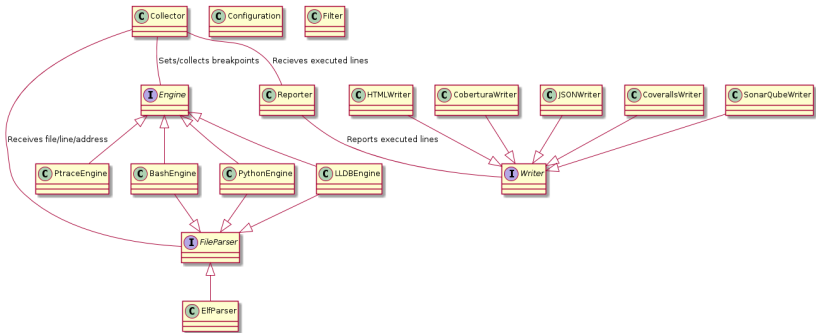
Example

```
auto child = fork();
if (child < 0)
    printf("Error\n");
else if (child == 0)
    printf("In child\n"); // Breakpoint here is bad news!
else
{
    printf("In parent, child %d\n", child);
    wait(&status);
}
```

So how does Python and Bash work?

- Bash has a **PS4** environment variable, which can be used to trace execution
- Python has a trace function which can be controlled via **sys.settrace**
- kcov needs to know what lines are “executable” and not, also in this setting
 - Easy in Python
 - Very tricky in Bash. Heredocs etc

Design



Questions and comments!

(Images from <http://www.falselogic.net/LetsPlay/SpaceQuest.html>)

Score: 202 of 202

Sound: on

